



CUDA ON ARM PLATFORM—线程层次

NVIDIA企业级开发者社区 何琨



加入 NVIDIA 开发者计划

获取最新版本软件、工具及开发信息


扫码了解详情



CUDA并行计算基础


- CUDA显存分配
- CUDA数据传输
- CUDA线程索引
- CUDA线程分配

我们要处理的问题




Element-wise

DAXPY



Local

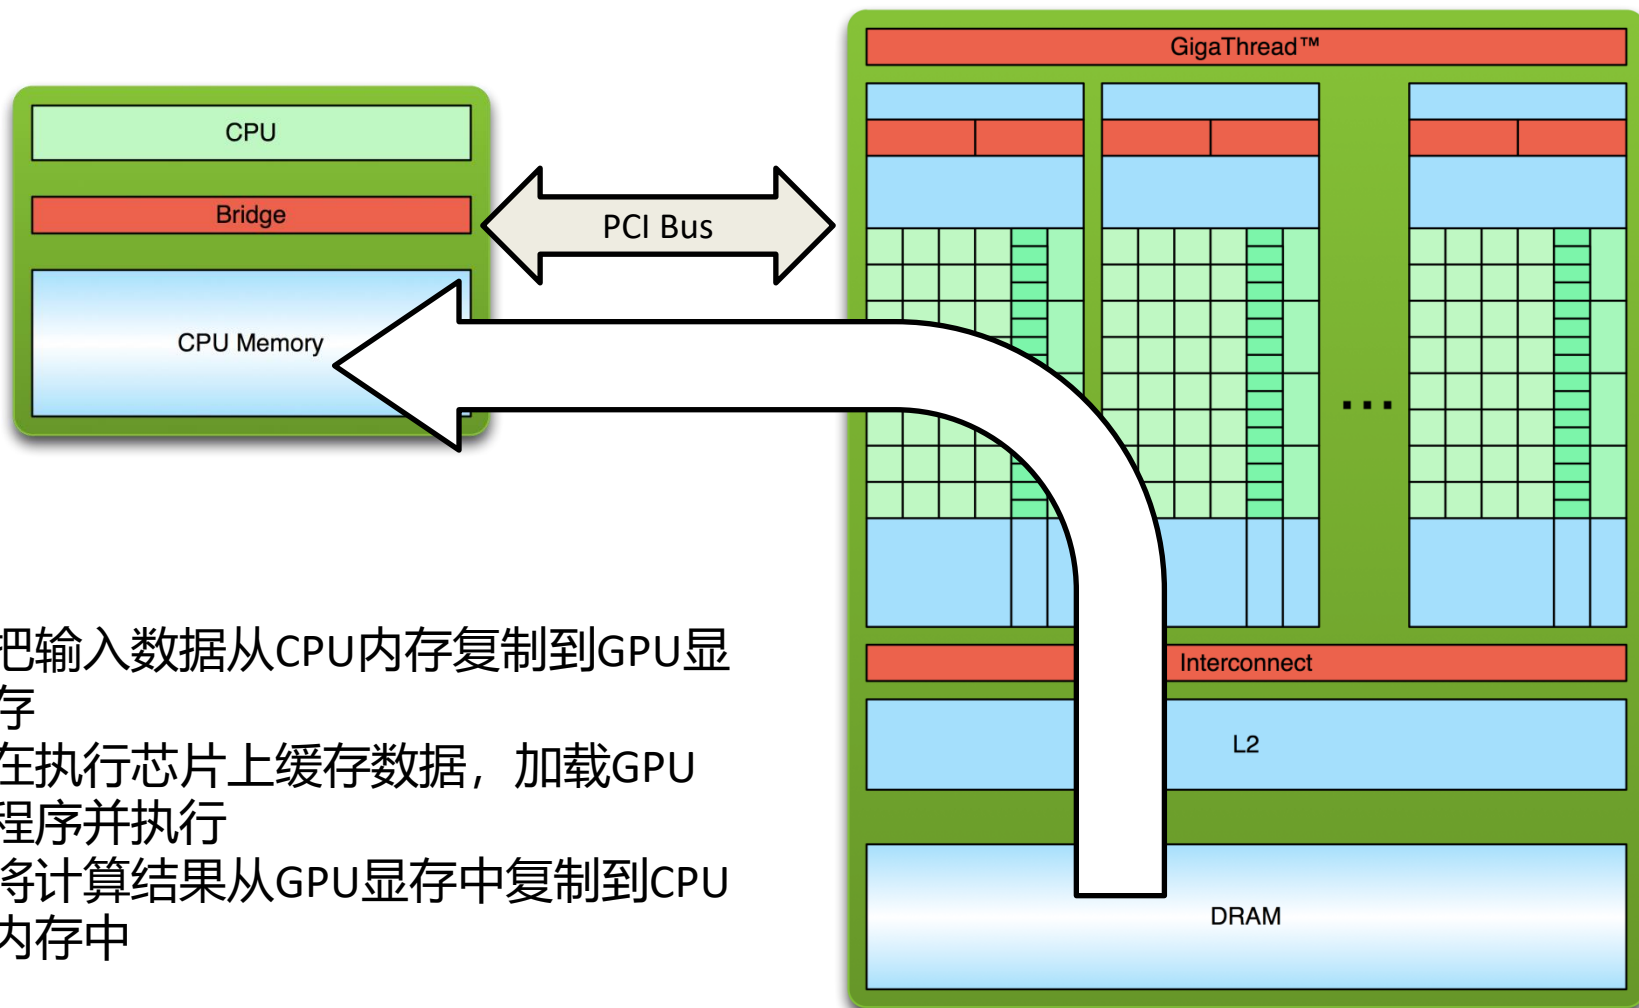
Convolution



All-to-All

Fourier Transform

CUDA程序的编写



MEMORY ALLOCATION

❖ `__host__ __device__ cudaError_t cudaMalloc(void** devPtr, size_t size)`

- **devPtr:**
 - - Pointer to allocated device memory
- **Size:**
 - - Requested allocation size in bytes

MEMORY COPY BETWEEN CPU AND GPU

❖ `cudaMemcpy(void *dst, const void *src, size_t count, cudaMemcpyKind kind)`


- **dst:** destination memory address
- **src:** source memory address
- **count:** size in bytes to copy
- **kind:** direction of the copy

❖ `cudaMemcpyKind`

- `cudaMemcpyHostToDevice`
- `cudaMemcpyDeviceToHost`
- `cudaMemcpyDeviceToDevice`
- `cudaMemcpyHostToHost`

CUDA的线程索引

- 如何确定线程执行地数据

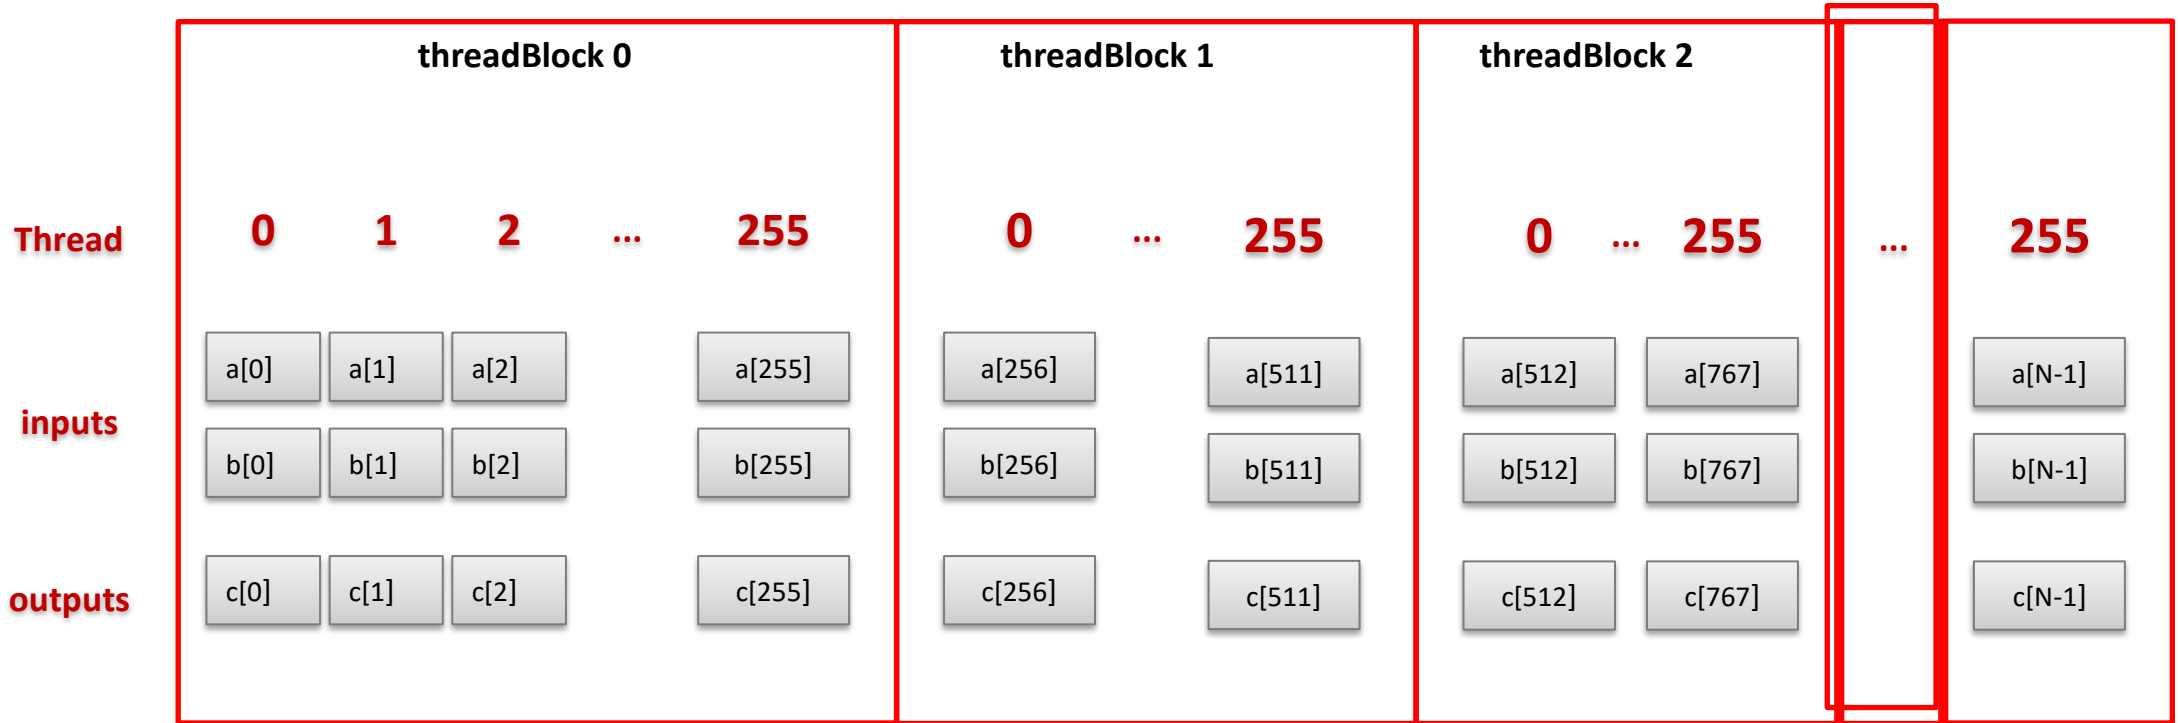


```
int index = threadIdx.x + blockIdx.x * blockDim.x;  
          =          5      +          2      * 8;  
          = 21;
```

PARALLELIZATION OF VECTORADD



PARALLELIZATION OF VECTORADD



work index $i = \text{threadIdx.x} + \text{blockIdx.x} * \text{blockDim.x};$

CUDA的线程索引

```
__global__ void add(const double *x, const double *y,  
double *z)  
{  
    const int n = blockDim.x * blockIdx.x + threadIdx.x;  
    z[n] = x[n] + y[n];  
}
```

每个线程都执行相同的命令

CUDA PROGRAMMING BY EXAMPLE


Case: Vector Add

❖ Parallelizable problem:

- $c = a + b$
- a, b, c are vectors of length N

❖ CPU implementation:

```
void main(){  
    int size = N * sizeof(int);  
    int *a, *b, *c;  
    a = (int *)malloc(size);  
    b = (int *)malloc(size);  
    c = (int *)malloc(size);  
    memset(c, 0, size);  
    init_rand_f(a, N);  
    init_rand_f(b, N);  
  
    vecAdd(N, a, b, c);  
}
```



```
void vecAdd (int n, int *a,  
             int *b, int *c)  
{  
    for(int i = 0; i < n; i++)  
    {  
        c[i] = a[i] + b[i];  
    }  
}
```

GPU CODE WORKFLOW

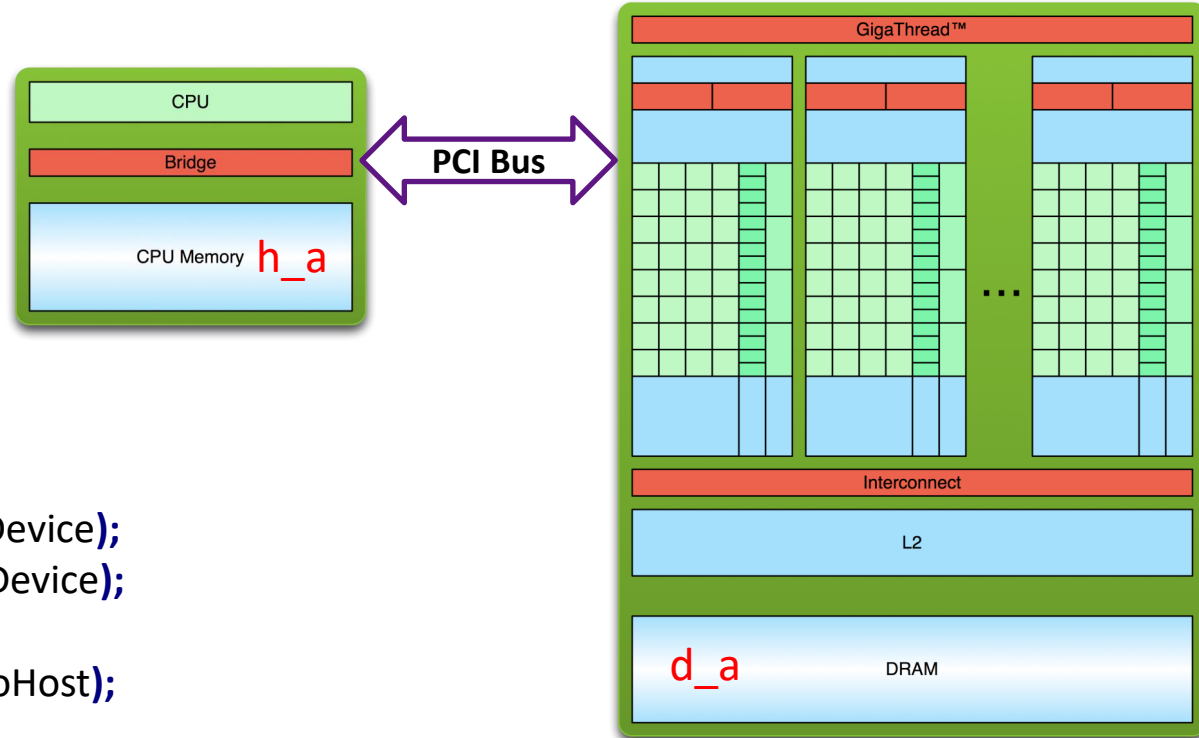
Allocate GPU Memories

```
int main(void) {  
    size_t size = N * sizeof(int);  
    int *h_a, *h_b; int *d_a, *d_b, *d_c;  
    h_a = (int *)malloc(size);  
    h_b = (int *)malloc(size);  
    ...
```

```
    cudaMalloc((void **)&d_a, size);  
    cudaMalloc((void **)&d_b, size);  
    cudaMalloc((void **)&d_c, size);
```

```
    cudaMemcpy(d_a, h_a, size, cudaMemcpyHostToDevice);  
    cudaMemcpy(d_b, h_b, size, cudaMemcpyHostToDevice);  
    vectorAdd<<<grid, block>>>>(d_a, d_b, d_c, N);  
    cudaMemcpy(h_c, d_c, size, cudaMemcpyDeviceToHost);
```

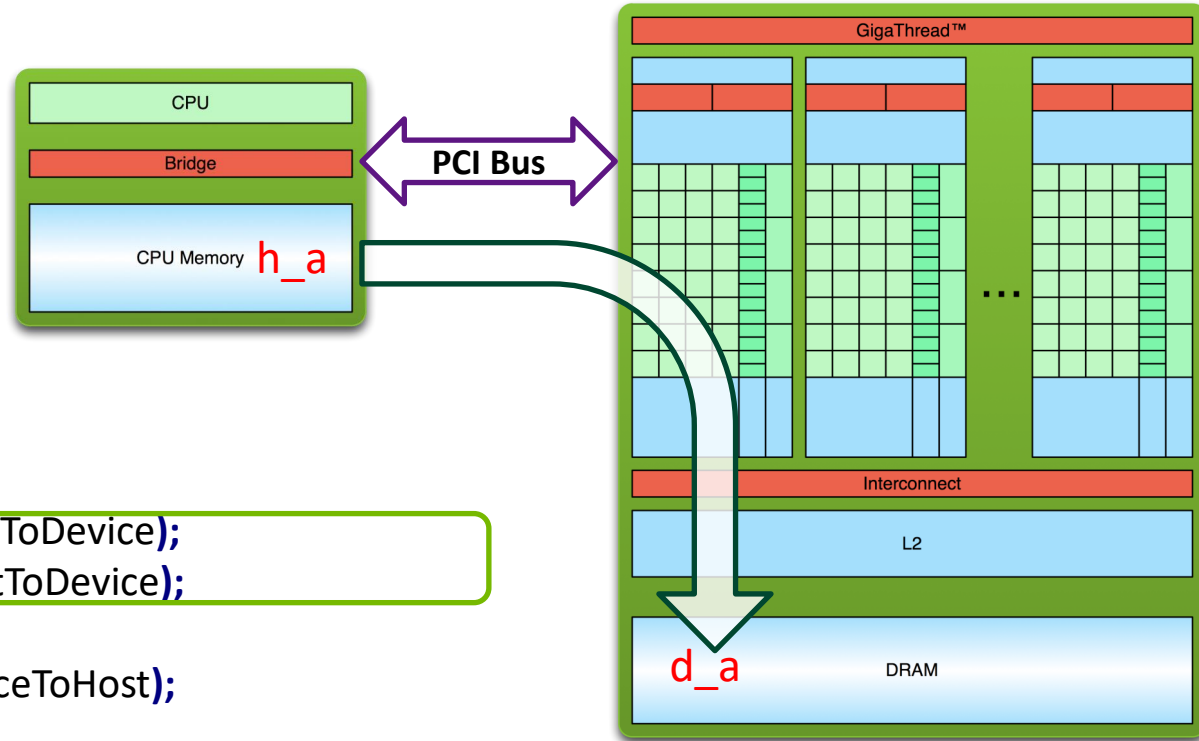
```
    cudaFree(d_a); cudaFree(d_b); cudaFree(d_c);  
    free(h_a); free(h_b);  
    return 0;}
```



GPU CODE WORKFLOW

Copy data from CPU to GPU

```
int main(void) {  
    size_t size = N * sizeof(int);  
    int *h_a, *h_b; int *d_a, *d_b, *d_c;  
    h_a = (int *)malloc(size);  
    h_b = (int *)malloc(size);  
    ...  
    cudaMalloc((void **)&d_a, size);  
    cudaMalloc((void **)&d_b, size);  
    cudaMalloc((void **)&d_c, size);  
  
    cudaMemcpy(d_a, h_a, size, cudaMemcpyHostToDevice);  
    cudaMemcpy(d_b, h_b, size, cudaMemcpyHostToDevice);  
    vectorAdd<<<grid, block>>>(d_a, d_b, d_c, N);  
    cudaMemcpy(h_c, d_c, size, cudaMemcpyDeviceToHost);  
  
    cudaFree(d_a); cudaFree(d_b); cudaFree(d_c);  
    free(h_a); free(h_b);  
    return 0;}
```

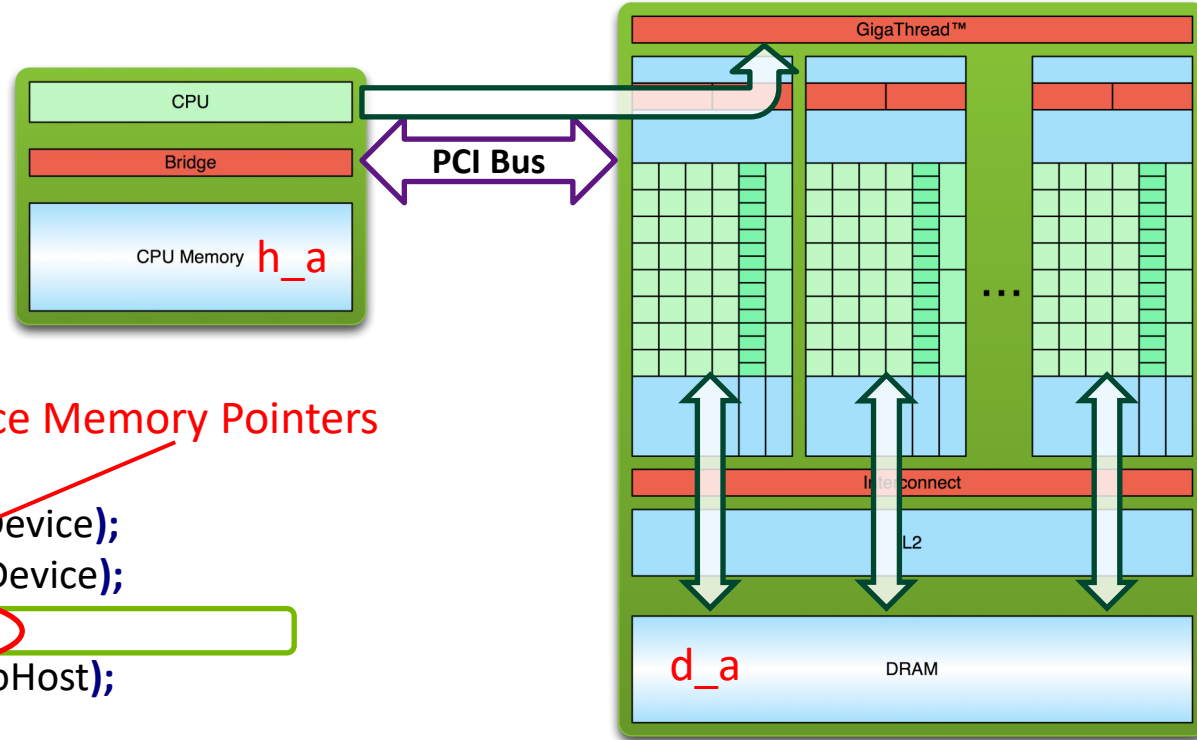


GPU CODE WORKFLOW

Invoke the CUDA Kernel

```
int main(void) {  
    size_t size = N * sizeof(int);  
    int *h_a, *h_b; int *d_a, *d_b, *d_c;  
    h_a = (int *)malloc(size);  
    h_b = (int *)malloc(size);  
    ...  
    cudaMalloc((void **)&d_a, size);  
    cudaMalloc((void **)&d_b, size);  
    cudaMalloc((void **)&d_c, size);  
  
    cudaMemcpy(d_a, h_a, size, cudaMemcpyHostToDevice);  
    cudaMemcpy(d_b, h_b, size, cudaMemcpyHostToDevice);  
    vectorAdd<<<grid, block>>>(d_a, d_b, d_c, N);  
    cudaMemcpy(h_c, d_c, size, cudaMemcpyDeviceToHost);  
  
    cudaFree(d_a); cudaFree(d_b); cudaFree(d_c);  
    free(h_a); free(h_b);  
    return 0;}
```

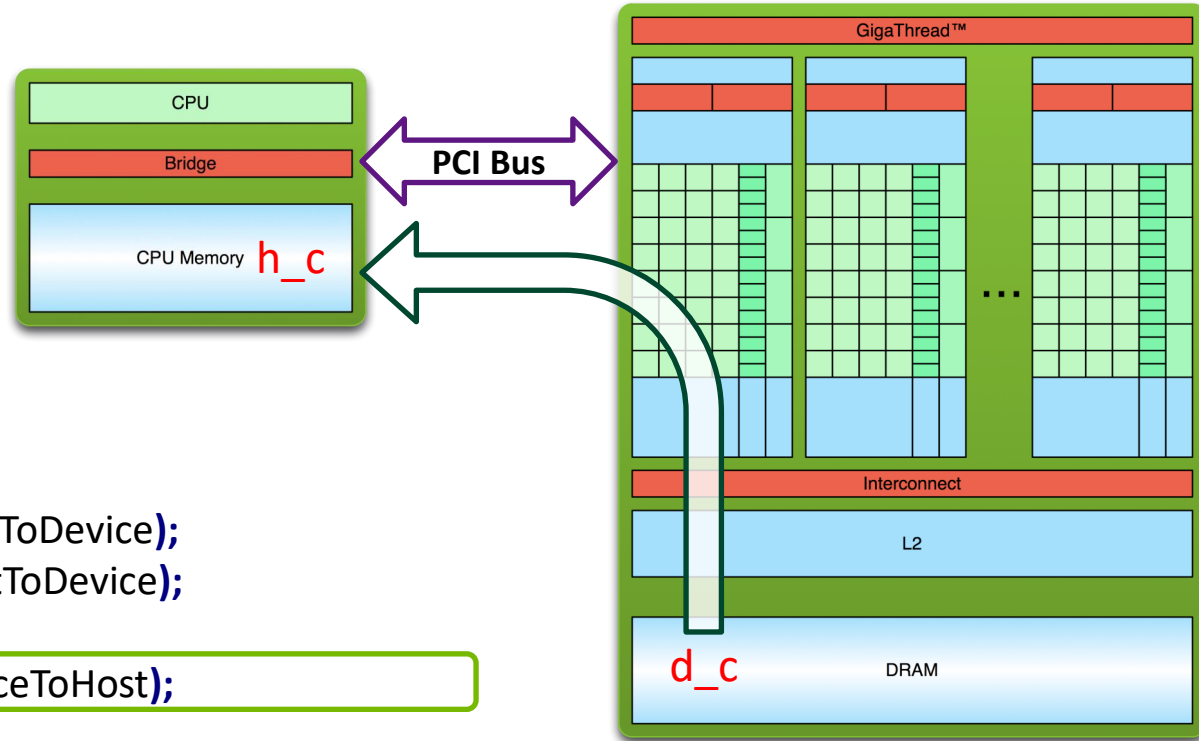
Device Memory Pointers



GPU CODE WORKFLOW

Copy result from GPU to CPU

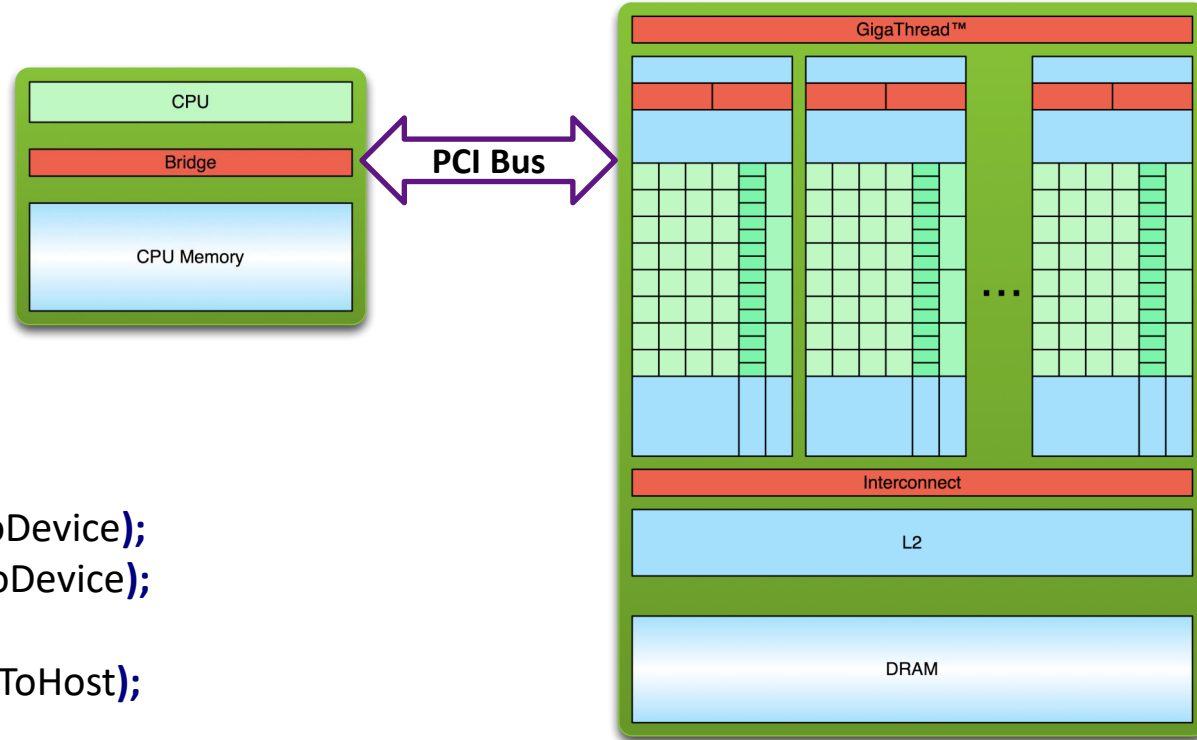
```
int main(void) {  
    size_t size = N * sizeof(int);  
    int *h_a, *h_b; int *d_a, *d_b, *d_c;  
    h_a = (int *)malloc(size);  
    h_b = (int *)malloc(size);  
    ...  
    cudaMalloc((void **)&d_a, size);  
    cudaMalloc((void **)&d_b, size);  
    cudaMalloc((void **)&d_c, size);  
  
    cudaMemcpy(d_a, h_a, size, cudaMemcpyHostToDevice);  
    cudaMemcpy(d_b, h_b, size, cudaMemcpyHostToDevice);  
    vectorAdd<<<grid, block>>>(d_a, d_b, d_c, N);  
    cudaMemcpy(h_c, d_c, size, cudaMemcpyDeviceToHost);  
  
    cudaFree(d_a); cudaFree(d_b); cudaFree(d_c);  
    free(h_a); free(h_b);  
    return 0;}
```



GPU CODE WORKFLOW

Release GPU Memories

```
int main(void) {  
    size_t size = N * sizeof(int);  
    int *h_a, *h_b; int *d_a, *d_b, *d_c;  
    h_a = (int *)malloc(size);  
    h_b = (int *)malloc(size);  
    ...  
    cudaMalloc((void **)&d_a, size);  
    cudaMalloc((void **)&d_b, size);  
    cudaMalloc((void **)&d_c, size);  
  
    cudaMemcpy(d_a, h_a, size, cudaMemcpyHostToDevice);  
    cudaMemcpy(d_b, h_b, size, cudaMemcpyHostToDevice);  
    vectorAdd<<<grid, block>>>(d_a, d_b, d_c, N);  
    cudaMemcpy(h_c, d_c, size, cudaMemcpyDeviceToHost);  
  
    cudaFree(d_a); cudaFree(d_b); cudaFree(d_c);  
    free(h_a); free(h_b);  
    return 0;}
```



CUDA的线程索引

如何设置Gridsize & Blocksize:

```
block_size = 128;  
grid_size = (N + block_size - 1) / block_size;
```

CUDA的线程分配

那么，我们的每个BLOCK可以申请多少个线程？

```
Total amount of shared memory per block:      49152 bytes
Total number of registers available per block: 65536
Warp size:                                     32
Maximum number of threads per multiprocessor:  2048
Maximum number of threads per block:           1024
Max dimension size of a thread block (x,y,z): (1024, 1024, 64)
Max dimension size of a grid size      (x,y,z): (2147483647, 65535, 65535)
```

CUDA的线程分配

那么，我们的每个BLOCK可以申请多少个线程？




CUDA的线程

那么，我们的每个BLOCK应该申请多少个线程？

CUDA的线程分配


WARP



❖ Warp is successive 32 threads in a block

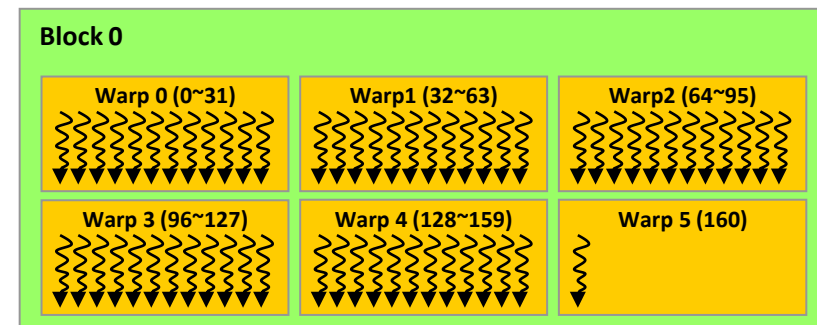
❖ E.g. blockDim = 160

— Automatically divided to 5 warps by GPU




❖ E.g. blockDim = 161

— If the blockDim is not the Multiple of 32
The rest of thread will occupy one more warp



CUDA的线程


那么，如果我们的数据过大，线程不够用怎么办？




0	4	8	12	16	20	24	28
1	5	9	13	17	21	25	29
2	6	10	14	18	22	26	30
3	7	11	15	19	23	27	31

CUDA的线程

那么，如果我们的数据过大，线程不够用怎么办？



0	4	8	12	16	20	24	28
1	5	9	13	17	21	25	29
2	6	10	14	18	22	26	30
3	7	11	15	19	23	27	31



红色框代表索引值为0的线程处理的数据？

CUDA的线程

那么，如果我们的数据过大，线程不够用怎么办？


```
__global__ add(const double *x, const double *y, double *z, int n)
{
    int index = blockDim.x * blockIdx.x + threadIdx.x;
    int stride = blockDim.x * gridDim.x;
    for(; index < n; index += stride)
        z[index] = x[index] + y[index];
}
```

CUDA的线程

那么，如果我们的数据过小，线程太多怎么办？

```
if( blockDim.x * blockIdx.x + threadIdx.x < count )
```

Sobel算子



$$\mathbf{G}_x = \begin{bmatrix} +1 & 0 & -1 \\ +2 & 0 & -2 \\ +1 & 0 & -1 \end{bmatrix} * \mathbf{A} \quad \text{and} \quad \mathbf{G}_y = \begin{bmatrix} +1 & +2 & +1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix} * \mathbf{A}$$

更多资源：



何琨-Ken

北京 密云



扫一扫上面的二维码图案，加我微信



加入 NVIDIA 开发者计划

获取最新版本软件、工具及开发信息

扫码了解详情



