




# CUDA ON ARM PLATFORM—矩阵乘 & 矩阵转置

NVIDIA企业级开发者社区 何琨



## 加入 NVIDIA 开发者计划

获取最新版本软件、工具及开发信息

扫码了解详情



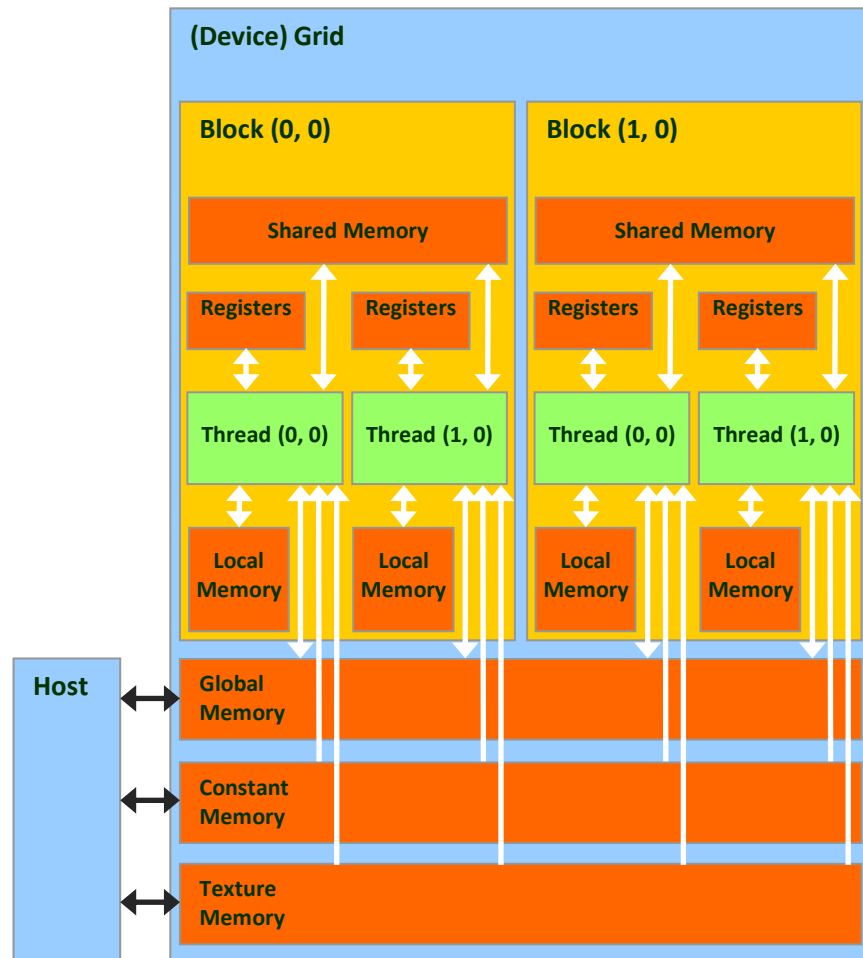
# CUDA并行计算基础

- GPU的存储单元
- 矩阵相乘
- 矩阵转置

# GPU的存储单元

Each thread can:

- R/W per-thread **registers**
  - R/W per-thread **local memory**
  - R/W per-block **shared memory**
  - R/W per-grid **global memory**
  - Read only per-grid **constant memory**
  - Read only per-grid **texture memory**
- The host **constant** **global** **texture**



# MEMORY ALLOCATION / RELEASE

CPU memory:

- `malloc()`
- `memset()`
- `free()`

GPU memory:

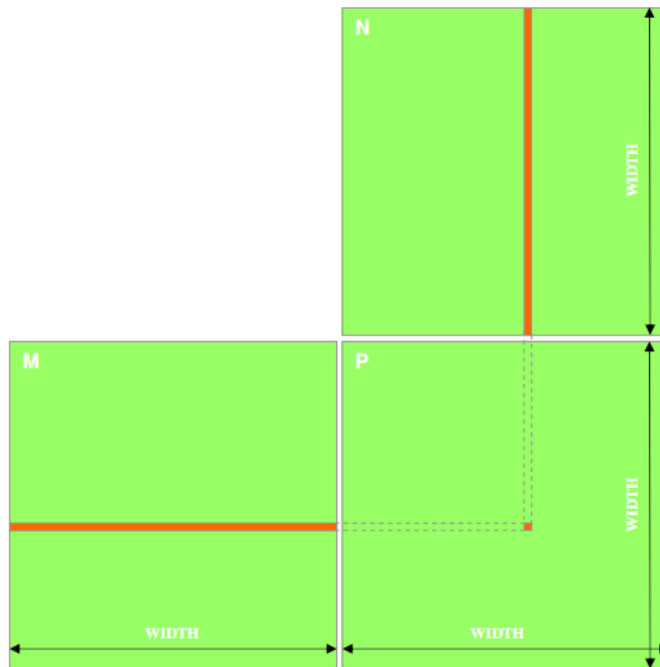
- `cudaMalloc()`
- `cudaMemset()`
- `cudaFree()`

# 矩阵相乘样例

```
void cpu_matrix_mult(int *h_m, int *h_n, int *h_result, int m,  
int n, int k) {  
    for (int i = 0; i < m; ++i)  
    {  
        for (int j = 0; j < k; ++j)  
        {  
            int tmp = 0.0;  
            for (int h = 0; h < n; ++h)  
            {  
                tmp += h_m[i * n + h] * h_n[h * k + j];  
            }  
            h_result[i * k + j] = tmp;  
        }  
    }  
}
```


$$P = M * N$$

假定 M and N 是方阵



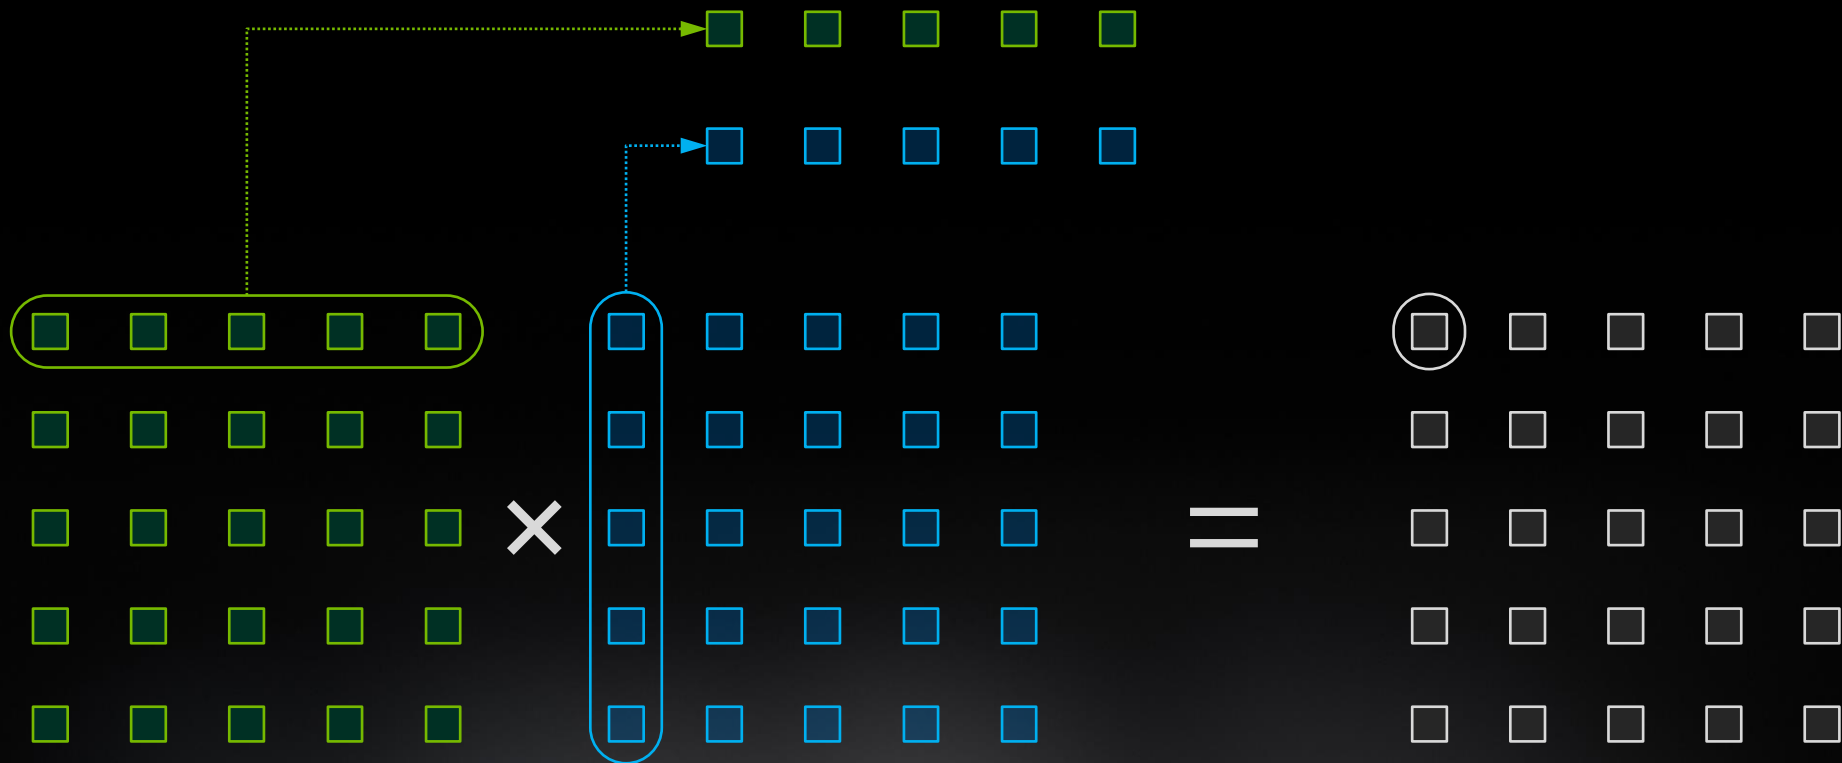
$$\begin{bmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \\ a_{31} & a_{32} & a_{33} \end{bmatrix} \times \begin{bmatrix} b_{11} & b_{12} & b_{13} \\ b_{21} & b_{22} & b_{23} \\ b_{31} & b_{32} & b_{33} \end{bmatrix} = \begin{bmatrix} a_{11} \times b_{11} + a_{12} \times b_{21} + a_{13} \times b_{31} & a_{11} \times b_{12} + a_{12} \times b_{22} + a_{13} \times b_{32} & a_{11} \times b_{13} + a_{12} \times b_{23} + a_{13} \times b_{33} \\ a_{21} \times b_{11} + a_{22} \times b_{21} + a_{23} \times b_{31} & a_{21} \times b_{12} + a_{22} \times b_{22} + a_{23} \times b_{32} & a_{21} \times b_{13} + a_{22} \times b_{23} + a_{23} \times b_{33} \\ a_{31} \times b_{11} + a_{32} \times b_{21} + a_{33} \times b_{31} & a_{31} \times b_{12} + a_{32} \times b_{22} + a_{33} \times b_{32} & a_{31} \times b_{13} + a_{32} \times b_{23} + a_{33} \times b_{33} \end{bmatrix}$$

# 矩阵乘示例

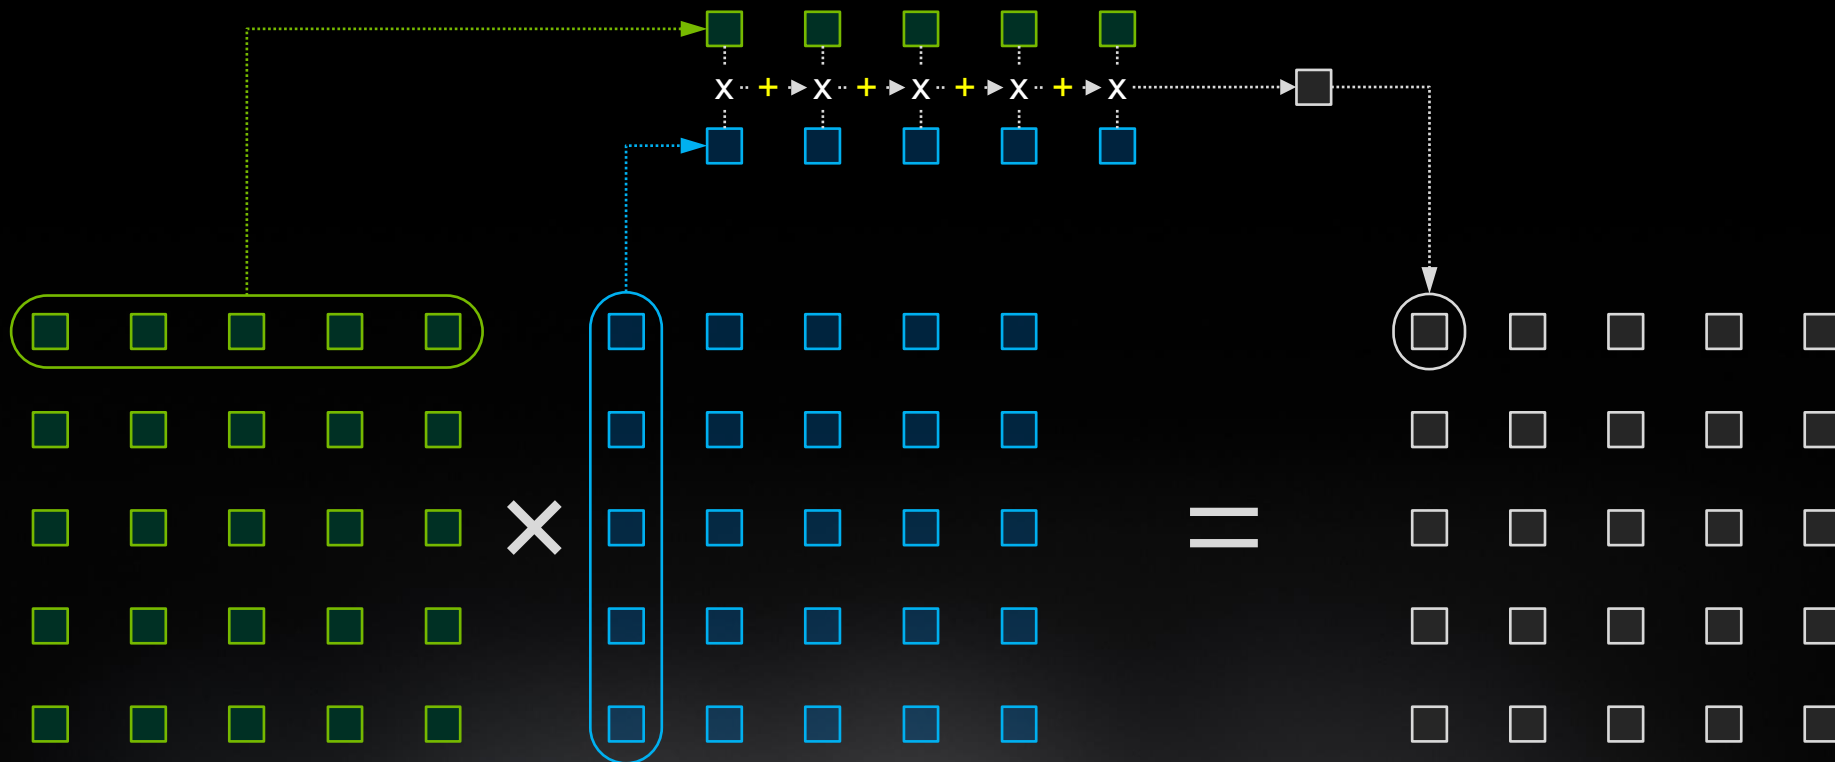




# 矩阵乘示例

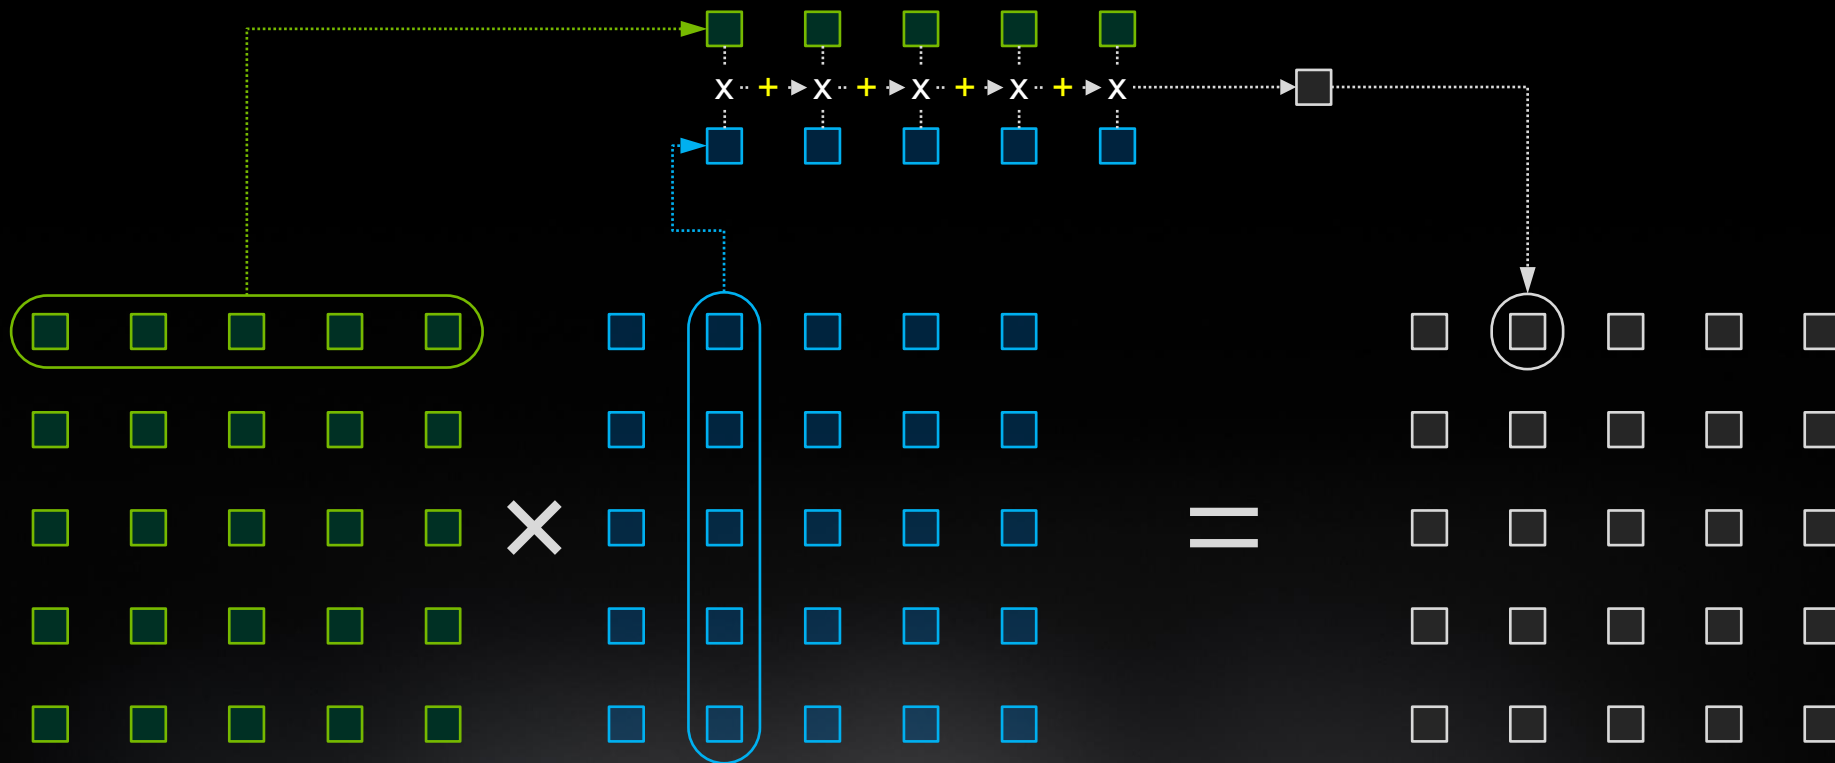


# 矩阵乘示例

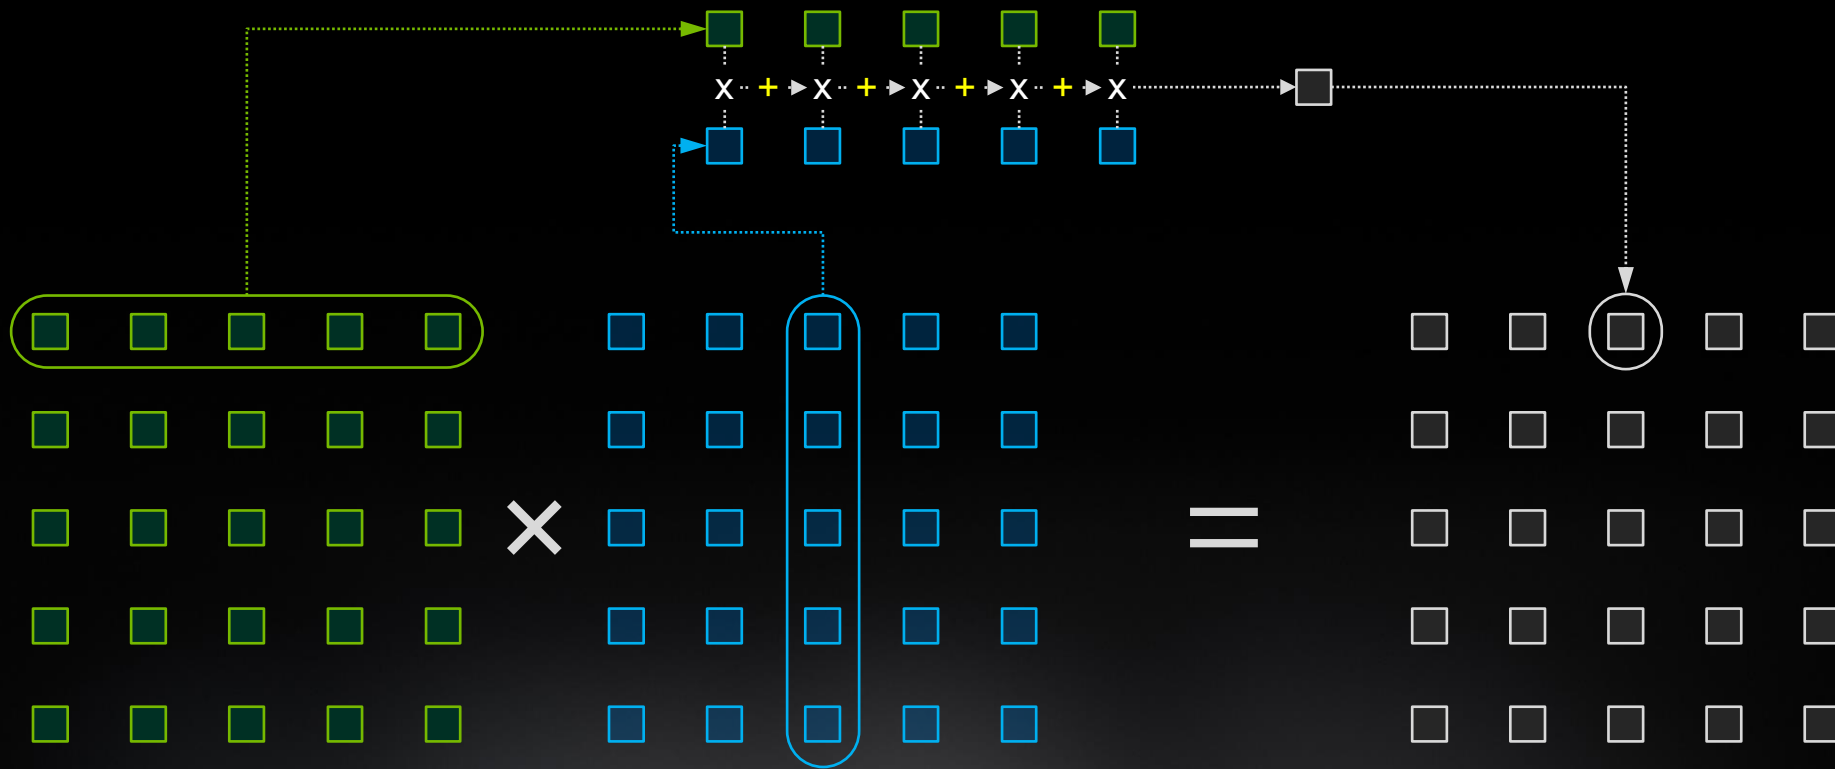




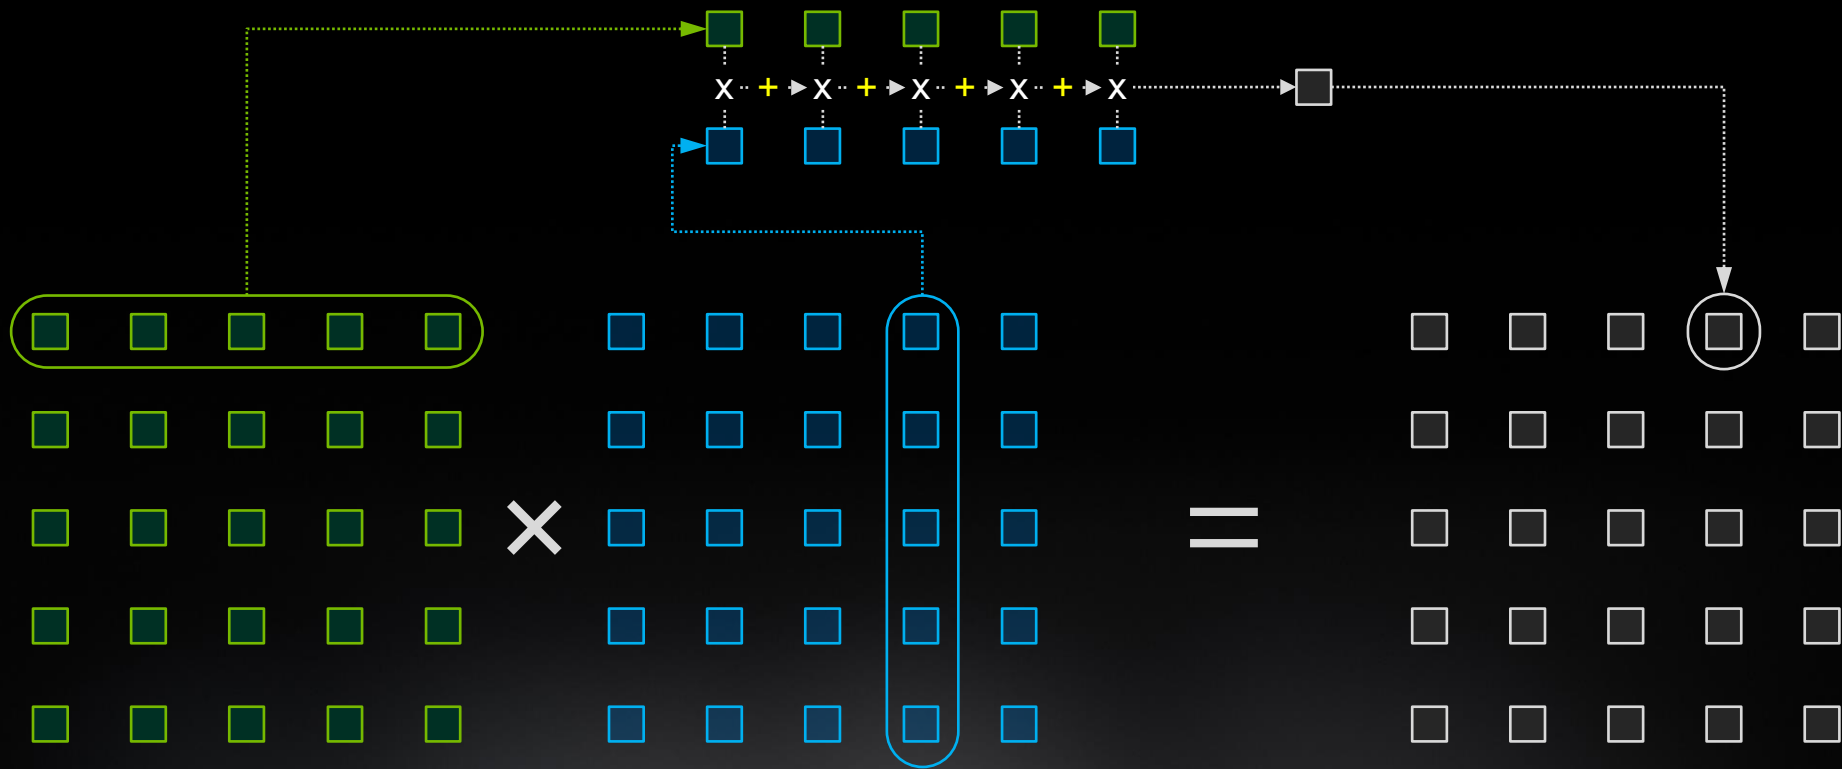
# 矩阵乘示例



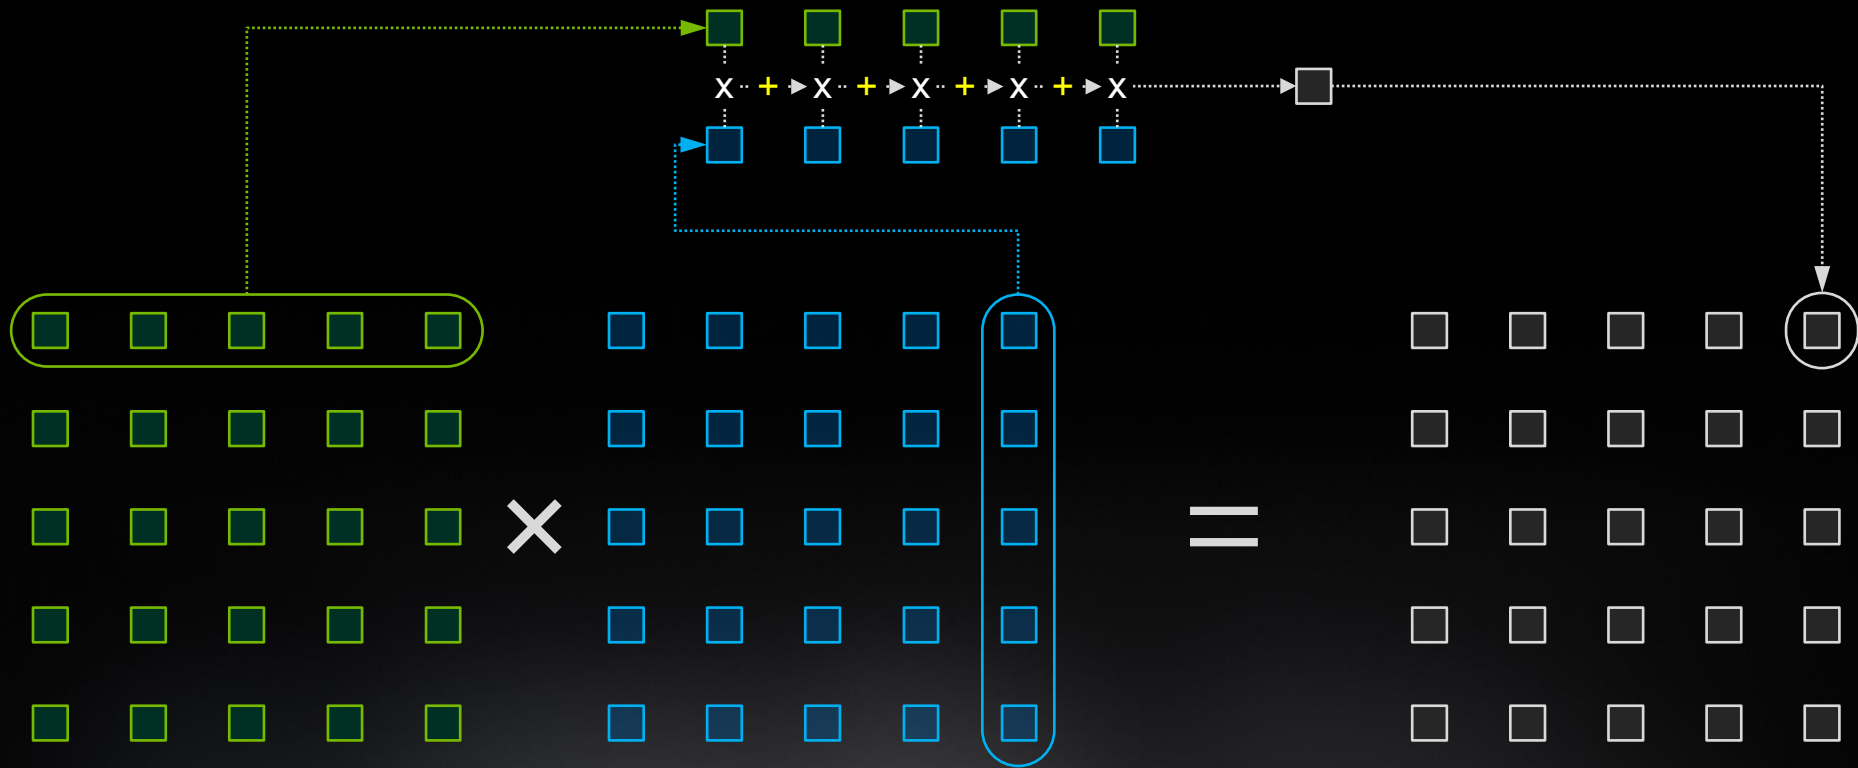
# 矩阵乘示例



# 矩阵乘示例



# 矩阵乘示例




# 矩阵相乘样例

```
void cpu_matrix_mult(int *h_m, int *h_n, int *h_result, int m,  
int n, int k) {  
    for (int i = 0; i < m; ++i)  
    {  
        for (int j = 0; j < k; ++j)  
        {  
            int tmp = 0.0;  
            for (int h = 0; h < n; ++h)  
            {  
                tmp += h_m[i * n + h] * h_n[h * k + j];  
            }  
            h_result[i * k + j] = tmp;  
        }  
    }  
}
```

$P = M * N$

假定 M and N 是方阵



Loop 1: j: 0

Loop 2: i: 0

Loop 3: k: 0

如果,  $M=N=1000$

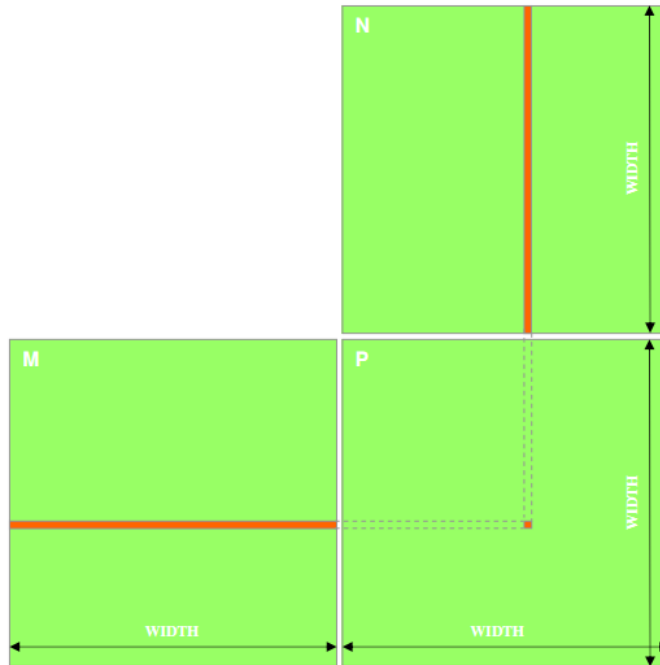
那么, 我们要做  $1000*1000*1000$  次:  $tmp += h\_m[i * n + h] * h\_n[h * k + j];$

# 矩阵相乘样例

```
void cpu_matrix_mult(int *h_m, int *h_n, int *h_result, int m,
int n, int k) {
    for (int i = 0; i < m; ++i)
    {
        for (int j = 0; j < k; ++j)
        {
            int tmp = 0.0;
            for (int h = 0; h < n; ++h)
            {
                tmp += h_m[i * n + h] * h_n[h * k + j];
            }
            h_result[i * k + j] = tmp;
        }
    }
}
```


$$P = M * N$$

假定 M and N 是方阵



那么，利用CUDA该怎么解决这个问题呢？  
空间换时间！

## 2D GRID AND 2D BLOCKS




								nx
0	1	2	3	4	5	6	7	Row 0
Block (0,0)				Block (1,0)				
8	9	10	11	12	13	14	15	Row 1
16	17	18	19	20	21	22	23	Row 2
Block (0,1)				Block (1,1)				
24	25	26	27	28	29	30	31	Row 3
32	33	34	35	36	37	38	39	Row 4
Block (0,2)				Block (1,2)				
40	41	42	43	44	45	46	47	Row 5
Col 0	Col 1	Col 2	Col 3	Col 4	Col 5	Col 6	Col 7	

```
__global__ void printThreadIndex(const int nx, const int ny) {
    int ix = threadIdx.x + blockIdx.x * blockDim.x;
    int iy = threadIdx.y + blockIdx.y * blockDim.y;
    unsigned int idx = iy*nx + ix;
    printf("thread_id (%d,%d) block_id (%d,%d) coordinate (%d, %d) "
           "global index %2d \n", threadIdx.x, threadIdx.y, blockIdx.x, blockIdx.y, ix, iy, idx);
}
```



# 2D GRID AND 2D BLOCKS



								$nx$	
	0	1	2	3	4	5	6	7	Row 0
	8	9	10	11	12	13	14	15	Row 1
	16	17	18	19	20	21	22	23	Row 2
	24	25	26	27	28	29	30	31	Row 3
	32	33	34	35	36	37	38	39	Row 4
	40	41	42	43	44	45	46	47	Row 5
	Col 0	Col 1	Col 2	Col 3	Col 4	Col 5	Col 6	Col 7	

Diagram illustrating a 2D grid of blocks with numerical indices. The grid is 8 rows by 8 columns. The horizontal axis is labeled  $nx$  and the vertical axis is labeled  $ny$ . The grid is divided into 2x2 blocks of 4x4 cells. The blocks are labeled as follows:

- Block (0,0) covers indices 0-15.
- Block (1,0) covers indices 16-31.
- Block (0,1) covers indices 32-47.
- Block (1,1) covers indices 48-63.

The cell at Row 5, Col 4 (index 44) is highlighted with a red box.

```
int index = (blockIdx.y*blockDim.y+threadIdx.y)*nx+ (blockIdx.x * blockDim.x + threadIdx.x);
```

```
Index == (2*2+1)*8+(1*4+0)
```

# 矩阵相乘样例

Grid											
0,0	1,0	2,0	3,0	0,0	1,0	2,0	3,0	0,0	1,0	2,0	3,0
0,1	1,1	2,1	3,1	0,0	1,1	2,1	3,1	0,1	1,1	2,1	3,1
0,2	1,2	2,2	3,2	0,2	1,2	2,2	3,2	0,2	1,2	2,2	3,2
0,0	1,1	2,0	3,0	0,0	1,0	2,0	3,0	0,0	1,0	2,0	3,0
0,0	1,1	2,1	3,1	0,0	1,1	2,1	3,1	0,1	1,1	2,1	3,1
0,0	1,2	2,2	3,2	0,0	1,2	2,2	3,2	0,2	1,2	2,2	3,2

每个颜色对应一个block  
该线程在grid中所有线程的索引为：

$\text{Thread\_x} = \text{blockIdx.x} * \text{blockDim.x} + \text{threadIdx.x} = 6$

$\text{Thread\_y} = \text{blockIdx.y} * \text{blockDim.y} + \text{threadIdx.y} = 3$

$\text{threadIdx.x} = 2$

$\text{threadIdx.y} = 0$

$\text{blockIdx.x} = 1$

$\text{blockIdx.y} = 1$

$\text{blockDim.x} = 4$

$\text{blockDim.y} = 3$

# 矩阵相乘样例

当把整个grid对应到一个矩阵的时候，它就像下面的样子：

0,0	1,0	2,0	3,0	4,0	5,0	6,0	7,0	8,0	9,0	10,0	11,0
0,1	1,1	2,1	3,1	4,1	5,1	6,1	7,1	8,1	9,1	10,1	11,1
0,2	1,2	2,2	3,2	4,2	5,2	6,2	7,2	8,2	9,2	10,2	11,2
0,3	1,3	2,3	3,3	4,3	5,3	6,3	7,3	8,3	9,3	10,3	11,3
0,4	1,4	2,4	3,4	4,4	5,4	6,4	7,4	8,4	9,4	10,4	11,4
0,5	1,5	2,5	3,5	4,5	5,5	6,5	7,5	8,5	9,5	10,5	11,5




而该线程在grid中所有线程的索引正好对应到矩阵的坐标。  
该线程在grid中所有线程的索引为：

$\text{Thread\_x} = \text{blockIdx.x} * \text{blockDim.x} + \text{threadIdx.x} = 6$

$\text{Thread\_y} = \text{blockIdx.y} * \text{blockDim.y} + \text{threadIdx.y} = 3$

# 矩阵相乘样例

- 一个线程grid计算Pd
  - 每个线程计算Pd的一个元素
- 每个线程
  - 读入矩阵Md的一行
  - 读入矩阵Nd的一列
  - 为每对Md和Nd元素执行一次乘法和加法



# 矩阵相乘样例

```
__global__ void gpu_matrix_mult(int *M, int *N, int *P, int
m_size, int n_size, int k_size)
{
    int row = blockIdx.y * blockDim.y + threadIdx.y;
    int col = blockIdx.x * blockDim.x + threadIdx.x;
    int sum = 0;
    if( col < k_size && row < m_size)
    {
        for(int i = 0; i < n; i++)
        {
            sum += M[row * n_size + i] * N[i * k_size + col];
        }
        P[row * k_size + col] = sum;
    }
}
```

计算出当前执行的线程在所有  
线程中的坐标



读取M矩阵的一行，N矩阵的  
一列，并做乘积累加



# 矩阵相乘样例

```
int main(int argc, char const *argv[])
{
    int m=100;
    int n=100;
    int k=100;

    int *h_a, *h_b, *h_c, *h_cc;
    cudaMallocHost((void **) &h_a, sizeof(int)*m*n);
    cudaMallocHost((void **) &h_b, sizeof(int)*n*k);
    cudaMallocHost((void **) &h_c, sizeof(int)*m*k);
    cudaMallocHost((void **) &h_cc, sizeof(int)*m*k);

    for (int i = 0; i < m; ++i) {
        for (int j = 0; j < n; ++j) {
            h_a[i * n + j] = rand() % 1024;
        }
    }

    for (int i = 0; i < n; ++i) {
        for (int j = 0; j < k; ++j) {
            h_b[i * k + j] = rand() % 1024;
        }
    }

    int *d_a, *d_b, *d_c;
    cudaMalloc((void **) &d_a, sizeof(int)*m*n);
    cudaMalloc((void **) &d_b, sizeof(int)*n*k);
    cudaMalloc((void **) &d_c, sizeof(int)*m*k);

    // copy matrix A and B from host to device memory
    cudaMemcpy(d_a, h_a, sizeof(int)*m*n, cudaMemcpyHostToDevice);
    cudaMemcpy(d_b, h_b, sizeof(int)*n*k, cudaMemcpyHostToDevice);

    unsigned int grid_rows = (m + BLOCK_SIZE - 1) / BLOCK_SIZE;
    unsigned int grid_cols = (k + BLOCK_SIZE - 1) / BLOCK_SIZE;
    dim3 dimGrid(grid_cols, grid_rows);
    dim3 dimBlock(BLOCK_SIZE, BLOCK_SIZE);

    gpu_matrix_mult<<<dimGrid, dimBlock>>>(d_a, d_b, d_c, m, n, k);

    cudaMemcpy(h_c, d_c, sizeof(int)*m*k, cudaMemcpyDeviceToHost);
    //cudaThreadSynchronize();
}
```

```
__global__ void gpu_matrix_mult(int *a,int *b, int *c, int m, int n, int k)
{
    int row = blockIdx.y * blockDim.y + threadIdx.y;
    int col = blockIdx.x * blockDim.x + threadIdx.x;
    int sum = 0;
    if( col < k && row < m)
    {
        for(int i = 0; i < n; i++)
        {
            sum += a[row * n + i] * b[i * k + col];
        }
        c[row * k + col] = sum;
    }
}
```

# 矩阵相乘样例

- 在算法实现中最主要的性能问题是什么?
- 主要的限制是什么?



# 矩阵转置示例

blockDim.x* blockIdx.x								threadIdx.x							
0, 0	0, 1	0, 2	0, 3	0, 4	0, 5	0, 6	0, 7	0, 8	0, 9	0, 10	0, 11	0, 12	0, 13	0, 14	0, 15
1, 0	1, 1	1, 2	1, 3	1, 4	1, 5	1, 6	1, 7	1, 8	1, 9	1, 10	1, 11	1, 12	1, 13	1, 14	1, 15
2, 0	2, 1	2, 2	2, 3	2, 4	2, 5	2, 6	2, 7	2, 8	2, 9	2, 10	2, 11	2, 12	2, 13	2, 14	2, 15
3, 0	3, 1	3, 2	3, 3	3, 4	3, 5	3, 6	3, 7	3, 8	3, 9	3, 10	3, 11	3, 12	3, 13	3, 14	3, 15
4, 0	4, 1	4, 2	4, 3	4, 4	4, 5	4, 6	4, 7	4, 8	4, 9	4, 10	4, 11	4, 12	4, 13	4, 14	4, 15
5, 0	5, 1	5, 2	5, 3	5, 4	5, 5	5, 6	5, 7	5, 8	5, 9	5, 10	5, 11	5, 12	5, 13	5, 14	5, 15
6, 0	6, 1	6, 2	6, 3	6, 4	6, 5	6, 6	6, 7	6, 8	6, 9	6, 10	6, 11	6, 12	6, 13	6, 14	6, 15
7, 0	7, 1	7, 2	7, 3	7, 4	7, 5	7, 6	7, 7	7, 8	7, 9	7, 10	7, 11	7, 12	7, 13	7, 14	7, 15
8, 0	8, 1	8, 2	8, 3	8, 4	8, 5	8, 6	8, 7	8, 8	8, 9	8, 10	8, 11	8, 12	8, 13	8, 14	8, 15
9, 0	9, 1	9, 2	9, 3	9, 4	9, 5	9, 6	9, 7	9, 8	9, 9	9, 10	9, 11	9, 12	9, 13	9, 14	9, 15
10, 0	10, 1	10, 2	10, 3	10, 4	10, 5	10, 6	10, 7	10, 8	10, 9	10, 10	10, 11	10, 12	10, 13	10, 14	10, 15
11, 0	11, 1	11, 2	11, 3	11, 4	11, 5	11, 6	11, 7	11, 8	11, 9	11, 10	11, 11	11, 12	11, 13	11, 14	11, 15
12, 0	12, 1	12, 2	12, 3	12, 4	12, 5	12, 6	12, 7	12, 8	12, 9	12, 10	12, 11	12, 12	12, 13	12, 14	12, 15
13, 0	13, 1	13, 2	13, 3	13, 4	13, 5	13, 6	13, 7	13, 8	13, 9	13, 10	13, 11	13, 12	13, 13	13, 14	13, 15
14, 0	14, 1	14, 2	14, 3	14, 4	14, 5	14, 6	14, 7	14, 8	14, 9	14, 10	14, 11	14, 12	14, 13	14, 14	14, 15
15, 0	15, 1	15, 2	15, 3	15, 4	15, 5	15, 6	15, 7	15, 8	15, 9	15, 10	15, 11	15, 12	15, 13	15, 14	15, 15

```
__global__ void gpu_transpose(int *in,int *out, int width)
{
    int y = blockIdx.y * blockDim.y + threadIdx.y;
    int x = blockIdx.x * blockDim.x + threadIdx.x;

    if( y < width && x < width)
    {
        out[x * width + y] = in[y * width + x];
    }
}
```

# 更多资源：

# <https://developer.nvidia-china.com>



何琨-Ken

北京 密云



# <https://www.nvidia.cn/developer/community-training/>

扫一扫上面的二维码图案，加我微信

